

Deep Learning in Practice



Anis Koubaa

The Vehicle Type Classification Project

In this notebook, we will present how to perform a classification of car brands. We will use different state of the art classifiers in Tensorflow 2.0 and Keras.

▼ Summary

- **Name:** Anis Koubaa
- **Date:** 20 September 2020
- **Use Case:** Vehicle Type
- **Algorithm:** MobileNetV2
- **Number of training images:** 603
- **Number of classes:** 7
- **Batch Size:** 64
- **Optimizer:** Adam
- **Learning Rate:** 0.0001
- **Loss Type:** CategoricalCrossentropy
- **Transfer Learning:** Yes | Imagenet

Comments: We obtained 100% on the validation accuracy on vehicle types, on validation dataset.

Let's get started.

We first need to load the requires libraries

```
# import the necessary packages
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import AveragePooling2D, GlobalAveragePooling2D, Batch
#from tensorflow.keras.applications import ResNet50
#from tensorflow.keras.applications import Xception
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.layers import AveragePooling2D
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Flatten
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import Input
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
```

```

.....,.....,.....,
headModel = Dense(number_of_classes, activation="softmax")(headModel)

# place the head FC model on top of the base model (this will become
# the actual model we will train)
model = Model(inputs=baseModel.input, outputs=headModel)

# loop over all layers in the base model and freeze them so they will
# *not* be updated during the first training process
for layer in baseModel.layers:
    layer.trainable = False

model.summary()

```



```

from tensorflow.keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical
from sklearn.preprocessing import LabelBinarizer
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from imutils import paths
import matplotlib.pyplot as plt
import numpy as np
import argparse
import cv2
import os
import sys
import tensorflow as tf
import h5py
import numpy as np
import sys

```

```
print(tf.__version__)
```

```
↳ 2.3.0
```

Then, we mount Google Drive to be able to access the files located on it

We now specify the path the dataset located on Google Drive

```

TYPE='type'
model_type='mobilenetv2'
user='anis'
iteration='2'

```

```
first_time_training=True
```

```

PROJECT_PATH='/content/drive/My Drive/udemy-deep-learning-in-practice/03-transfer-l
print('PROJECT_PATH: ',PROJECT_PATH)
HDF5_DATASET_PATH=PROJECT_PATH+'datasets/vehicle-type-dataset-SIZE224-train-dev-tes
print('HDF5_DATASET_PATH: ', HDF5_DATASET_PATH)
ACCURACY_LOSS_OUPUT_FILE=PROJECT_PATH+'log/'+model_type+'/'+model_type+'-by-'+TYPE+
TARGET_CLASSIFICATION_MODEL=PROJECT_PATH+'trained-models/'+model_type+'/'+model_type+'-by-'+TYPE+
print('TARGET_CLASSIFICATION_MODEL: ',TARGET_CLASSIFICATION_MODEL)
CHECKPOINT_PATH = PROJECT_PATH+'checkpoints/'+model_type+'/'+model_type+'-by-'+TYPE+'-'+model_t
print('CHECKPOINT_PATH: ',CHECKPOINT_PATH)
LOGFILE_PATH=PROJECT_PATH+'log/'+model_type+'/'+model_type+'-by-'+TYPE+'-training-l
print('LOGFILE_PATH: ',LOGFILE_PATH)

```

```

↳ PROJECT_PATH: /content/drive/My Drive/udemy-deep-learning-in-practice/03-trar
HDF5_DATASET_PATH: /content/drive/My Drive/udemy-deep-learning-in-practice/03
TARGET_CLASSIFICATION_MODEL: /content/drive/My Drive/udemy-deep-learning-in-practice/03-trar
CHECKPOINT_PATH: /content/drive/My Drive/udemy-deep-learning-in-practice/03-trar
LOGFILE_PATH: /content/drive/My Drive/udemy-deep-learning-in-practice/03-trar

```

```
from google.colab import drive
drive.mount('/content/drive')
```

☞ Drive already mounted at /content/drive; to attempt to forcibly remount, call

```
sys.path.append(PROJECT_PATH)
import anis_koubaa_udemy_computer_vision_lib
from anis_koubaa_udemy_computer_vision_lib import *
```

▼ Load the Dataset

```
def load_dataset_from_hdf5_file(hdf_file_path):
    hf = h5py.File(hdf_file_path, "r")
    trainX= np.array(hf["trainX"]).astype("f8")
    ascii_train_labels = np.array(hf["trainLabels"]).astype("S65")
    trainY=np.array(hf["trainY"]).astype("int")

    devX= np.array(hf["devX"]).astype("f8")
    ascii_dev_labels = np.array(hf["devLabels"]).astype("S65")
    devY=np.array(hf["devY"]).astype("int")

    testX= np.array(hf["testX"]).astype("f8")
    ascii_test_labels = np.array(hf["testLabels"]).astype("S65")
    testY=np.array(hf["testY"]).astype("int")

    trainLabels = np.array([n.decode('unicode_escape') for n in ascii_train_labels])
    devLabels = np.array([n.decode('unicode_escape') for n in ascii_dev_labels])
    testLabels = np.array([n.decode('unicode_escape') for n in ascii_test_labels])

    print("trainX.shape: ",trainX.shape)
    print("trainY.shape: ",trainY.shape)
    print("trainLabels.shape: ",trainLabels.shape)
    print("devX.shape: ",devX.shape)
    print("devY.shape: ",devY.shape)
    print("devLabels.shape: ",devLabels.shape)
    print("testX.shape: ",testX.shape)
    print("testY.shape: ",testY.shape)
    print("testLabels.shape: ",testLabels.shape)

    return trainX, trainY, trainLabels, devX,devY,devLabels,testX,testY,testLabels

trainX, trainY, trainLabels, devX,devY,devLabels,testX,testY,testLabels=load_datase
```

☞

```
trainX.shape: (603, 224, 224, 3)
trainY.shape: (603, 7)
trainLabels.shape: (603,)
devX.shape: (75, 224, 224, 3)
devY.shape: (75, 7)
```

```
IMAGE_SIZE=trainX.shape[1]
print(IMAGE_SIZE)
```

```
224
```

```
number_of_classes=np.unique(trainLabels).size
```

Dataset Visualization

```
anis_koubaa_udemy_computer_vision_lib.plot_sample_from_dataset(trainX, trainLabels,
```

```
224,
```




motorcycle-motorbike-chopper



motorcycle-motorbike-chopper



motorcycle-bicycle-racing



car-s



motorcycle-bicycle-racing



car-suv-alltypes




car-suv-alltypes



car-t

```
anis_koubaa_udemy_computer_vision_lib.plot_sample_from_dataset(devX, devLabels,rows
```

```
224,
```



```
anis_koubaa_udemy_computer_vision_lib.plot_sample_from_dataset(testX, testLabels,ro
```



car-bus-alltypes



car-sedan-alltypes



car-bus-alltypes



car-se

▼ Training Configuration

Data Augmentation

```
# initialize the training data augmentation object
trainAug = ImageDataGenerator(
    rotation_range=20)
#fill_mode="nearest")
#brightness_range=[0.2,1.0])
#horizontal_flip=True)

# load the network, ensuring the head FC layer sets are left
# off
if (first_time_training==True):
    print('training for first time')
    baseModel = MobileNetV2(weights="imagenet", include_top=False, input_shape=(IMA

# training for first time
```

```
# construct the head of the model that will be placed on top of the
# the base model
headModel = baseModel.output
headModel = AveragePooling2D(pool_size=(4, 4))(headModel)
headModel = Flatten(name="flatten")(headModel)
headModel = Dense(128, activation="relu")(headModel)
headModel = BatchNormalization()(headModel)
headModel = Dense(64, activation="relu")(headModel)
headModel = Dropout(0.5)(headModel)
headModel = BatchNormalization()(headModel)
```

block_13_depthwise_BN (BatchNorm	(None, 1, 1, 516)	2304	block_13_dept
block_13_depthwise_relu (ReLU)	(None, 7, 7, 576)	0	block_13_dept
block_13_project (Conv2D)	(None, 7, 7, 160)	92160	block_13_dept
block_13_project_BN (BatchNorma	(None, 7, 7, 160)	640	block_13_proj
block_14_expand (Conv2D)	(None, 7, 7, 960)	153600	block_13_proj
block_14_expand_BN (BatchNormal	(None, 7, 7, 960)	3840	block_14_exp
block_14_expand_relu (ReLU)	(None, 7, 7, 960)	0	block_14_exp
block_14_depthwise (DepthwiseCo	(None, 7, 7, 960)	8640	block_14_exp
block_14_depthwise_BN (BatchNor	(None, 7, 7, 960)	3840	block_14_dept
block_14_depthwise_relu (ReLU)	(None, 7, 7, 960)	0	block_14_dept
block_14_project (Conv2D)	(None, 7, 7, 160)	153600	block_14_dept
block_14_project_BN (BatchNorma	(None, 7, 7, 160)	640	block_14_proj
block_14_add (Add)	(None, 7, 7, 160)	0	block_13_proj block_14_proj
block_15_expand (Conv2D)	(None, 7, 7, 960)	153600	block_14_add[
block_15_expand_BN (BatchNormal	(None, 7, 7, 960)	3840	block_15_exp
block_15_expand_relu (ReLU)	(None, 7, 7, 960)	0	block_15_exp
block_15_depthwise (DepthwiseCo	(None, 7, 7, 960)	8640	block_15_exp
block_15_depthwise_BN (BatchNor	(None, 7, 7, 960)	3840	block_15_dept
block_15_depthwise_relu (ReLU)	(None, 7, 7, 960)	0	block_15_dept
block_15_project (Conv2D)	(None, 7, 7, 160)	153600	block_15_dept
block_15_project_BN (BatchNorma	(None, 7, 7, 160)	640	block_15_proj
block_15_add (Add)	(None, 7, 7, 160)	0	block_14_add[block_15_proj
block_16_expand (Conv2D)	(None, 7, 7, 960)	153600	block_15_add[
block_16_expand_BN (BatchNormal	(None, 7, 7, 960)	3840	block_16_exp
block_16_expand_relu (ReLU)	(None, 7, 7, 960)	0	block_16_exp
block_16_depthwise (DepthwiseCo	(None, 7, 7, 960)	8640	block_16_exp
block_16_depthwise_BN (BatchNor	(None, 7, 7, 960)	3840	block_16_dept
block_16_depthwise_relu (ReLU)	(None, 7, 7, 960)	0	block_16_dept
block_16_project (Conv2D)	(None, 7, 7, 320)	307200	block_16_dept
block_16_project_BN (BatchNorma	(None, 7, 7, 320)	1280	block_16_proj

Conv_1 (Conv2D)	(None, 7, 7, 1280)	409600	block_16_proj
Conv_1_bn (BatchNormalization)	(None, 7, 7, 1280)	5120	Conv_1[0][0]
out_relu (ReLU)	(None, 7, 7, 1280)	0	Conv_1_bn[0][0]
average_pooling2d (AveragePooling2D)	(None, 1, 1, 1280)	0	out_relu[0][0]
flatten (Flatten)	(None, 1280)	0	average_pooling2d[0][0]
dense (Dense)	(None, 128)	163968	flatten[0][0]
batch_normalization (BatchNormalization)	(None, 128)	512	dense[0][0]
dense_1 (Dense)	(None, 64)	8256	batch_normalization[0][0]
dropout (Dropout)	(None, 64)	0	dense_1[0][0]
batch_normalization_1 (BatchNormalization)	(None, 64)	256	dropout[0][0]
dense_2 (Dense)	(None, 7)	455	batch_normalization_1[0][0]
=====			
Total params: 2,431,431			
Trainable params: 173,063			
Non-trainable params: 2,258,368			

#disable this instruction if you train for the first time